

Interactive exercise week #9b

Liping Wu 300-958-061 11/9/2020

In this exercise we will do the following:

- Explore a dataset
- Process & clean up and visualize the dataset
- Build a logistic regression model

Pre-requisites:

- 1- Install Anaconda
- 2- We will be using a lot of Public datasets these datasets are available at <https://goo.gl/zjS4C6> under a folder named "Datasets for Predictive Modelling with Python", the datasets are organized in the order of the text book chapters: Python: Advanced Predictive Analytics, [chapter # 6](#) files are required

Steps for exploring and building a logistic regression model:

- 1- Open your spider IDE
- 2- Load the 'Bank.csv' file into a dataframe name *the dataframe data_firstname_b where first name is your first name* carry out the following activities:
 - a. Display the column names
 - b. Display the shape of the data frame i.e number of rows and number of columns
 - c. Display the main statistics of the data
 - d. Display the types of columns
 - e. Display the first five records

Following is the code, *make sure you update the path to the correct path where you placed the files and update the data frame name correctly:*

```
import pandas as pd
import os
path = " D:/CentennialWu/2020Fall/COMP309Data/Assignments/Lab07Wk09"
filename = 'Bank.csv'
fullpath = os.path.join(path,filename)
data_liping_b = pd.read_csv(fullpath,sep=';')

print(data_liping_b.columns.values)
```

```
print(data_liping_b.shape)
print(data_liping_b.describe())
print(data_liping_b.dtypes)
print(data_liping_b.head(5))
```

```
...: print(data_liping_b.head(5))
['age' 'job' 'marital' 'education' 'default' 'housing' 'loan' 'contact'
 'month' 'day_of_week' 'duration' 'campaign' 'pdays' 'previous' 'poutcome'
 'emp.var.rate' 'cons.price.idx' 'cons.conf.idx' 'euribor3m' 'nr.employed'
 'y']
(4119, 21)
count  4119.000000  age      duration  ...  euribor3m  nr.employed
mean    40.113620   256.788055  ...    3.621356   5166.481695
std     10.313362   254.703736  ...    1.733591    73.667904
min     18.000000    0.000000  ...    0.635000   4963.600000
25%     32.000000   103.000000  ...    1.334000   5099.100000
50%     38.000000   181.000000  ...    4.857000   5191.000000
75%     47.000000   317.000000  ...    4.961000   5228.100000
max     88.000000  3643.000000  ...    5.045000   5228.100000

[8 rows x 10 columns]
age                int64
job                object
marital            object
education          object
default            object
housing            object
loan               object
contact            object
month              object
day_of_week        object
duration           int64
campaign           int64
pdays             int64
previous           int64
poutcome           object
emp.var.rate       float64
cons.price.idx     float64
cons.conf.idx      float64
euribor3m          float64
nr.employed        float64
y                  object
dtype: object
   age  job      marital  ...  euribor3m  nr.employed  y
0   30  blue-collar  married  ...    1.313     5099.1  no
1   39  services    single  ...    4.855     5191.0  no
2   25  services  married  ...    4.962     5228.1  no
3   38  services  married  ...    4.959     5228.1  no
4   47   admin.  married  ...    4.191     5195.8  no
```

- 3- Explore, Check & Process the data, we need to carry out several processing steps (clean up steps) before building the model as follows:
 - a. 1-change the y column from object to integer, this is the class or label column
 - b. Reduce the categories of the education column
 - c. Check the values of who purchased the deposit account
 - d. Check the average of all the numeric columns
 - e. Check the mean of all numeric columns grouped by education
 - f. Plot a histogram showing purchase by education category

- g. Draw a stacked bar chart of the marital status and the purchase of term deposit to see whether this can be a good predictor of the outcome
- h. Plot the bar chart for the Frequency of Purchase against each day of the week to see whether this can be a good predictor of the outcome.
- i. Repeat step h for the month
- j. Plot a histogram of the age distribution

Following is the code:

Explore, Check & Process the data

1-change the y column from object to integer

```
print(data_liping_b)
```

```
data_liping_b['y']=(data_liping_b['y']=='yes').astype(int)
```

```
print(data_liping_b)
```

```
In [6]: print(data_liping_b)
...: data_liping_b['y']=(data_liping_b['y']=='yes').astype(int)
...: print(data_liping_b)
```

	age	job	marital	...	euribor3m	nr.employed	y
0	30	blue-collar	married	...	1.313	5099.1	no
1	39	services	single	...	4.855	5191.0	no
2	25	services	married	...	4.962	5228.1	no
3	38	services	married	...	4.959	5228.1	no
4	47	admin.	married	...	4.191	5195.8	no
...
4114	30	admin.	married	...	4.958	5228.1	no
4115	39	admin.	married	...	4.959	5228.1	no
4116	27	student	single	...	1.354	5099.1	no
4117	58	admin.	married	...	4.966	5228.1	no
4118	34	management	single	...	4.120	5195.8	no

```
[4119 rows x 21 columns]
```

	age	job	marital	...	euribor3m	nr.employed	y
0	30	blue-collar	married	...	1.313	5099.1	0
1	39	services	single	...	4.855	5191.0	0
2	25	services	married	...	4.962	5228.1	0
3	38	services	married	...	4.959	5228.1	0
4	47	admin.	married	...	4.191	5195.8	0
...
4114	30	admin.	married	...	4.958	5228.1	0
4115	39	admin.	married	...	4.959	5228.1	0
4116	27	student	single	...	1.354	5099.1	0
4117	58	admin.	married	...	4.966	5228.1	0
4118	34	management	single	...	4.120	5195.8	0

```
[4119 rows x 21 columns]
```

#2- Reduce the categories of the education column

```
print(data_liping_b['education'].unique())
```

```
In [14]: print(data_liping_b['education'].unique())
['basic.9y' 'high.school' 'university.degree' 'professional.course'
'basic.6y' 'basic.4y' 'unknown' 'illiterate']
```

```
import numpy as np
```

```
data_liping_b['education']=np.where(data_liping_b['education']=='basic.9y', 'Basic',
data_liping_b['education'])
```

```
data_liping_b['education']=np.where(data_liping_b['education']=='basic.6y', 'Basic',
data_liping_b['education'])
```

```

data_liping_b['education']=np.where(data_liping_b['education']=='basic.4y', 'Basic',
data_liping_b['education'])
data_liping_b['education']=np.where(data_liping_b['education']=='university.degree', 'University
Degree', data_liping_b['education'])
data_liping_b['education']=np.where(data_liping_b['education']=='professional.course',
'Professional Course', data_liping_b['education'])
data_liping_b['education']=np.where(data_liping_b['education']=='high.school', 'High School',
data_liping_b['education'])
data_liping_b['education']=np.where(data_liping_b['education']=='illiterate', 'Illiterate',
data_liping_b['education'])
data_liping_b['education']=np.where(data_liping_b['education']=='unknown', 'Unknown',
data_liping_b['education'])

```

After cleaning

```
print(data_liping_b['education'].unique())
```

```

...: print(data_liping_b['education'].unique())
['Basic' 'High School' 'University Degree' 'Professional Course' 'Unknown'
'Illiterate']

```

#Check the values of who purchased the deposit account

```
print(data_liping_b['y'].value_counts())
```

```

In [16]: print(data_liping_b['y'].value_counts())
0      3668
1       451
Name: y, dtype: int64

```

#Check the average of all the numeric columns

```
pd.set_option('display.max_columns',100)
```

```
print(data_liping_b.groupby('y').mean())
```

```

In [18]: pd.set_option('display.max_columns',100)
...: print(data_liping_b.groupby('y').mean())

```

	age	duration	campaign	pdays	previous	emp.var.rate	\
y							
0	39.895311	219.40976	2.605780	982.763086	0.141767	0.240185	
1	41.889135	560.78714	1.980044	778.722838	0.585366	-1.177384	

	cons.price.idx	cons.conf.idx	euribor3m	nr.employed
y				
0	93.599677	-40.586723	3.802826	5175.502072
1	93.417268	-39.786475	2.145448	5093.118625

#Check the mean of all numeric columns grouped by education

```
print(data_liping_b.groupby('education').mean())
```

```
In [19]: print(data_liping_b.groupby('education').mean())
```

	age	duration	campaign	pdays	previous	\
education						
Basic	42.337124	253.898457	2.429732	978.815597	0.149472	
High School	38.097720	258.534202	2.630836	958.022801	0.206298	
Illiterate	42.000000	146.000000	4.000000	999.000000	0.000000	
Professional Course	40.207477	278.816822	2.512150	958.211215	0.194393	
University Degree	39.017405	247.707278	2.583070	947.900316	0.207278	
Unknown	42.826347	267.281437	2.538922	939.700599	0.263473	

	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m	\
education					
Basic	0.237368	93.658600	-41.120552	3.775701	
High School	-0.002497	93.564314	-40.995765	3.511732	
Illiterate	-2.900000	92.201000	-31.400000	0.834000	
Professional Course	0.163925	93.599630	-40.127664	3.701426	
University Degree	-0.009731	93.499109	-39.830063	3.547132	
Unknown	-0.074251	93.637455	-39.487425	3.410174	

	nr.employed	y
education		
Basic	5174.133144	0.079610
High School	5163.212595	0.105320
Illiterate	5076.200000	0.000000
Professional Course	5167.595140	0.121495
University Degree	5163.023180	0.130538
Unknown	5151.260479	0.155689

#Plot a histogram showing purchase by education category

import matplotlib.pyplot as plt

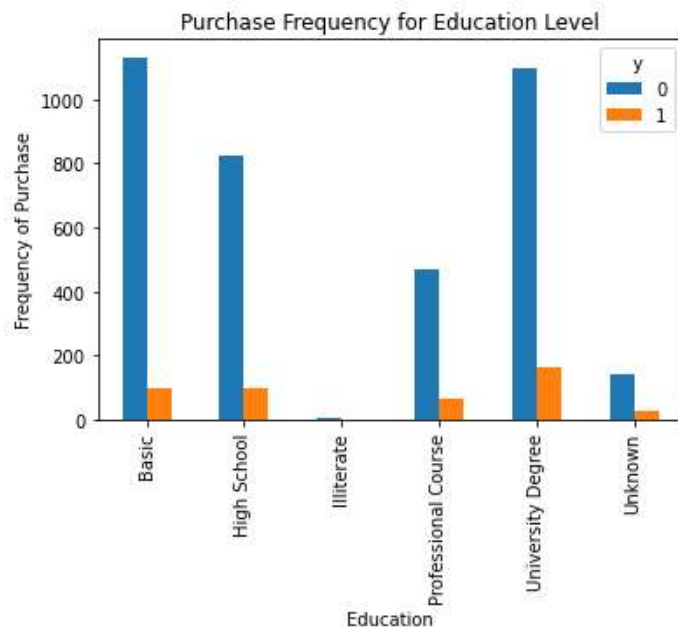
pd.crosstab(data_liping_b.education,data_liping_b.y)

pd.crosstab(data_liping_b.education,data_liping_b.y).plot(kind='bar')

plt.title('Purchase Frequency for Education Level')


plt.xlabel('Education')

plt.ylabel('Frequency of Purchase')



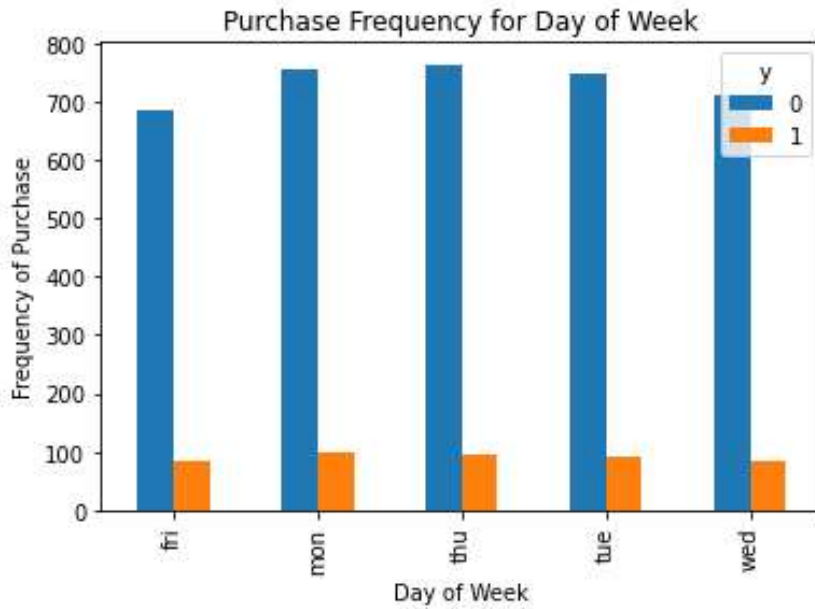
#draw a stacked bar chart of the marital status and the purchase of term deposit to see whether this can be a good predictor of the outcome

```
table=pd.crosstab(data_liping_b.marital,data_liping_b.y)
table.div(table.sum(1).astype(float), axis=0).plot(kind='bar', stacked=True)
plt.title('Stacked Bar Chart of Marital Status vs Purchase')
plt.xlabel('Marital Status')
plt.ylabel('Proportion of Customers')
```




#plot the bar chart for the Frequency of Purchase against each day of the week to see whether this can be a good predictor of the outcome

```
pd.crosstab(data_liping_b.day_of_week,data_liping_b.y).plot(kind='bar')
plt.title('Purchase Frequency for Day of Week')
plt.xlabel('Day of Week')
plt.ylabel('Frequency of Purchase')
```


#Repeat for the month

```
pd.crosstab(data_liping_b.month,data_liping_b.y).plot(kind='bar')
plt.title('Purchase Frequency for Day of Week')
plt.xlabel('Month')
plt.ylabel('Frequency of Purchase')
```



#Plot a histogram of the age distribution

```
data_liping_b.age.hist()
plt.title('Histogram of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
```



- 4- Deal with the categorical variables, as follows:
- Create the dummy variables , use a loop
 - Remove the original columns
 - Prepare the data for the model build as X (inputs, predictor) and Y(output, predicted)

Following is the code.

```
cat_vars=['job','marital','education','default','housing','loan','contact','month','day_of_week','pou  
tcome']  
for var in cat_vars:  
    cat_list='var'+ '_' +var  
    print(cat_list)  
    cat_list = pd.get_dummies(data_liping_b[var], prefix=var)  
    print(cat_list)
```


cat_list.describe

```
In [131]: cat_list.describe
Out[131]:
<bound method NDFrame.describe of      poutcome_failure  poutcome_nonexistent  poutcome_success
0                0                1                0
1                0                1                0
2                0                1                0
3                0                1                0
4                0                1                0
...
4114            0                1                0
4115            0                1                0
4116            1                0                0
4117            0                1                0
4118            0                1                0

[4119 rows x 3 columns]>
```

print(data_liping_b)

```
      poutcome  emp.var.rate  cons.price.idx  cons.conf.idx  euribor3m  \
0  nonexistent      -1.8      92.893      -46.2      1.313
1  nonexistent       1.1      93.994      -36.4      4.855
2  nonexistent       1.4      94.465      -41.8      4.962
3  nonexistent       1.4      94.465      -41.8      4.959
4  nonexistent      -0.1      93.200      -42.0      4.191
...
4114  nonexistent       1.4      93.918      -42.7      4.958
4115  nonexistent       1.4      93.918      -42.7      4.959
4116    failure      -1.8      92.893      -46.2      1.354
4117  nonexistent       1.4      93.444      -36.1      4.966
4118  nonexistent      -0.1      93.200      -42.0      4.120

      nr.employed  y
0      5099.1  0
1      5191.0  0
2      5228.1  0
3      5228.1  0
4      5195.8  0
...
4114      5228.1  0
4115      5228.1  0
4116      5099.1  0
4117      5228.1  0
4118      5195.8  0

[4119 rows x 21 columns]
```

#data_liping_b1=data_liping_b.join(cat_list,how= 'inner') need to fix by join

data_liping_b = pd.merge(data_liping_b, cat_list,left_index=True, right_index=True, how='outer')

to observe the datachange

data_liping_b.columns.values

data_liping_b.describe()

data_liping_b.head()

data_liping_b.describe

```

      nr.employed  y  poutcome_failure  poutcome_nonexistent  poutcome_success
0      5099.1  0      0      1      0
1      5191.0  0      0      1      0
2      5228.1  0      0      1      0
3      5228.1  0      0      1      0
4      5195.8  0      0      1      0
...
4114      5228.1  0      0      1      0
4115      5228.1  0      0      1      0
4116      5099.1  0      1      0      0
4117      5228.1  0      0      1      0
4118      5195.8  0      0      1      0

[4119 rows x 24 columns]>

```

2- Remove the original columns

```
cat_vars=['job','marital','education','default','housing','loan','contact','month','day_of_week','poutcome']
```

```
data_liping_b_vars=data_liping_b.columns.values.tolist()
```

```
to_keep=[i for i in data_liping_b_vars if i not in cat_vars]
```

```
data_liping_b_final=data_liping_b[to_keep]
```

```
data_liping_b_final.columns.values
```

```

In [135]: data_liping_b_final.columns.values
Out[135]:
array(['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate',
      'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed', 'y',
      'poutcome_failure', 'poutcome_nonexistent', 'poutcome_success'],
      dtype=object)

```

```
data_liping_b_final.describe
```

```

In [136]: data_liping_b_final.describe
Out[136]:
<bound method NDFrame.describe of      age  duration  campaign  pdays  previous  emp.var.rate  cons.price.idx  \
0      30      487         2    999         0        -1.8      92.893
1      39      346         4    999         0         1.1      93.994
2      25      227         1    999         0         1.4      94.465
3      38       17         3    999         0         1.4      94.465
4      47       58         1    999         0        -0.1      93.200
...    ...    ...    ...    ...    ...    ...    ...
4114   30       53         1    999         0         1.4      93.918
4115   39      219         1    999         0         1.4      93.918
4116   27       64         2    999         1        -1.8      92.893
4117   58      528         1    999         0         1.4      93.444
4118   34      175         1    999         0        -0.1      93.200

      cons.conf.idx  euribor3m  nr.employed  y  poutcome_failure  \
0      -46.2         1.313      5099.1  0         0
1      -36.4         4.855      5191.0  0         0
2      -41.8         4.962      5228.1  0         0
3      -41.8         4.959      5228.1  0         0
4      -42.0         4.191      5195.8  0         0
...    ...    ...    ...    ...    ...
4114   -42.7         4.958      5228.1  0         0
4115   -42.7         4.959      5228.1  0         0
4116   -46.2         1.354      5099.1  0         1
4117   -36.1         4.966      5228.1  0         0
4118   -42.0         4.120      5195.8  0         0

      poutcome_nonexistent  poutcome_success
0              1              0
1              1              0
2              1              0
3              1              0
4              1              0
...    ...    ...
4114              1              0
4115              1              0
4116              0              0
4117              1              0
4118              1              0

[4119 rows x 14 columns]>

```

3- Prepare the data for the model build as X (inputs, predictor) and Y(output, predicted)

data_liping_b_final_vars=data_liping_b_final.columns.values.tolist()

Y=['y']

X=[i for i in data_liping_b_final_vars if i not in Y]

type(Y)

type(X)

```

In [137]: data_liping_b_final_vars=data_liping_b_final.columns.values.tolist()

In [138]: Y=['y']
...: X=[i for i in data_liping_b_final_vars if i not in Y ]

In [139]: type(Y)
...: type(X)
Out[139]: list

In [140]: data_liping_b_final_vars=data_liping_b_final.columns.values.tolist()
...: Y=['y']
...: X=[i for i in data_liping_b_final_vars if i not in Y ]
...: type(Y)
...: type(X)
Out[140]: list

```

- 5- Carryout feature selection and update the data, as follows:
- Carry out feature selection using the RFE module from sklearn.model_selection to select only 12 feature
 - Update X and Y to reflect only 12 features

Following is the code:

```
from sklearn import datasets
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
rfe = RFE(model, 12)
rfe = rfe.fit(data_liping_b_final[X],data_liping_b_final[Y] )
print(rfe.support_)
print(rfe.ranking_)
for r in zip(rfe.support_, rfe.ranking_):
    print(r)
```

```
In [144]: print(rfe.support_)
[ True  True  True False  True  True  True  True  True  True  True  True
 True]

In [145]: print(rfe.ranking_)
[1 1 1 2 1 1 1 1 1 1 1 1]

In [146]: for r in zip(rfe.support_, rfe.ranking_):
...:     print(r)
(True, 1)
(True, 1)
(True, 1)
(False, 2)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
(True, 1)
```

#2- Update X and Y with selected features

```
cols=['previous', 'euribor3m', 'poutcome_success', 'poutcome_failure']
data_liping_b_final.columns.values
X=data_liping_b_final[cols]
Y=data_liping_b_final['y']
type(Y)
type(X)
```

```

In [154]: cols=['previous', 'euribor3m', 'poutcome_success', 'poutcome_failure']
...: data_liping_b_final.columns.values
Out[154]:
array(['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate',
      'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed', 'y',
      'poutcome_failure', 'poutcome_nonexistent', 'poutcome_success'],
      dtype=object)

In [155]: X=data_liping_b_final[cols]

In [156]: Y=data_liping_b_final['y']

In [157]: type(Y)
...: type(X)
Out[157]: pandas.core.frame.DataFrame

```

- 6- Build the logistic regression model as follows:
 - a. Split the data into 70%training and 30% for testing
 - b. Build the model using “sklearn linear_model.LogisticRegression”
 - c. Fit the training data
 - d. Validate the parameters and check model accuracy

Following is the code:

#1- split the data into 70%training and 30% for testing, note added the solver to avoid warnings

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, random_state=0)
```

```
X_train.describe # [2883 rows x 4 columns]>
```

```
Y_train.describe # Name: y, Length: 2883, dtype: int32>
```

```
X_test.describe
```

```
Y_test.describe
```

```

In [166]: X_train.describe
Out[166]:
<bound method NDFrame.describe of      previous  euribor3m  poutcome_success  poutcome_failure
4070      0      4.021      0      0
1240      0      4.856      0      0
349       0      4.963      0      0
3706      0      4.866      0      0
4043      0      4.857      0      0
...      ...      ...      ...      ...
1033      0      4.961      0      0
3264      0      4.961      0      0
1653      0      0.879      0      0
2607      0      4.962      0      0
2732      0      4.966      0      0

[2883 rows x 4 columns]>

In [167]: Y_train.describe
Out[167]:
<bound method NDFrame.describe of 4070  0
1240  0
349   0
3706  0
4043  0
..
1033  0
3264  0
1653  1
2607  0
2732  1
Name: y, Length: 2883, dtype: int32>

```

```

In [168]: X_test.describe
Out[168]:
<bound method NDFrame.describe of      previous  euribor3m  poutcome_success  poutcome_failure
3754      0      4.961      0      0
45      0      1.327      0      0
2774      0      4.865      0      0
1170      0      4.958      0      0
4042      0      4.120      0      0
...      ...      ...      ...      ...
1728      0      4.958      0      0
1460      0      4.960      0      0
1363      0      0.884      0      0
1898      0      4.857      0      0
3519      1      1.726      1      0

[1236 rows x 4 columns]>

In [169]: Y_test.describe
Out[169]:
<bound method NDFrame.describe of 3754      0
45      0
2774      0
1170      0
4042      0
...      ..
1728      0
1460      0
1363      0
1898      0
3519      1
Name: y, Length: 1236, dtype: int32>

```

2-Let us build the model and validate the parameters

```
from sklearn import linear_model
```

```
from sklearn import metrics
```

```
clf1 = linear_model.LogisticRegression(solver='lbfgs')
```

```
clf1.fit(X_train, Y_train)
```

#3- Run the test data against the new model

```
probs = clf1.predict_proba(X_test)
```

```
print(probs)
```

```
predicted = clf1.predict(X_test)
```

```
print (predicted)
```

```

In [175]: print(probs)
[[0.95653402 0.04346598]
 [0.8027594  0.1972406 ]
 [0.95464216 0.04535784]
 ...
 [0.76814958 0.23185042]
 [0.95448101 0.04551899]
 [0.47625222 0.52374778]]

```

#4-Check model accuracy

```
print (metrics.accuracy_score(Y_test, predicted))
```

```

In [177]: print (metrics.accuracy_score(Y_test, predicted))
0.9037216828478964

```

7- To avoid sampling bias run cross validation for 10 times, as follows:

- Use the cross_val_score module from sklearn.model_selection and set the parameters
- Save the results of each run in scores
- Produce the mean

Following is the code:

```
from sklearn.model_selection import cross_val_score
scores = cross_val_score(linear_model.LogisticRegression(solver='lbfgs'), X, Y,
scoring='accuracy', cv=10)
print (scores)
print (scores.mean())
```

```
In [178]: from sklearn.model_selection import cross_val_score
...: scores = cross_val_score(linear_model.LogisticRegression(solver='lbfgs'), X, Y,
...: scoring='accuracy', cv=10)
...: print (scores)
...: print (scores.mean())
[0.91990291 0.90048544 0.8907767  0.89805825 0.90533981 0.90533981
 0.89320388 0.90048544 0.89320388 0.90754258]
0.901433869558028
```

8- Generate the confusion matrix as follows:

- Prepare two arrays one for the predicted values Y_P and one for actual values Y_A of the test. For the predicted use a threshold of 0.05, this means if the probability is higher than 0.05 the model will classify the instance as 1 and if it is lower than 0.05 it will be classified as 0.
- Use the `confusion_matrix` option from the `sklearn.metrics` module to generate the matrix

Following is the code:

```
prob=probs[:,1]
prob_df=pd.DataFrame(prob)
prob_df['predict']=np.where(prob_df[0]>=0.05,1,0)
import numpy as np
Y_A=Y_test.values
Y_P = np.array(prob_df['predict'])
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(Y_A, Y_P)
print (confusion_matrix)
```

```
In [179]: prob=probs[:,1]
...: prob_df=pd.DataFrame(prob)
...: prob_df['predict']=np.where(prob_df[0]>=0.05,1,0)
...: import numpy as np
...: Y_A=Y_test.values
...: Y_P = np.array(prob_df['predict'])
...: from sklearn.metrics import confusion_matrix
...: confusion_matrix = confusion_matrix(Y_A, Y_P)
...: print (confusion_matrix)
[[715 399]
 [ 35  87]]
```